


# Computing the Stratified Semantics of Logic Programs over Big Data through Mass Parallelization

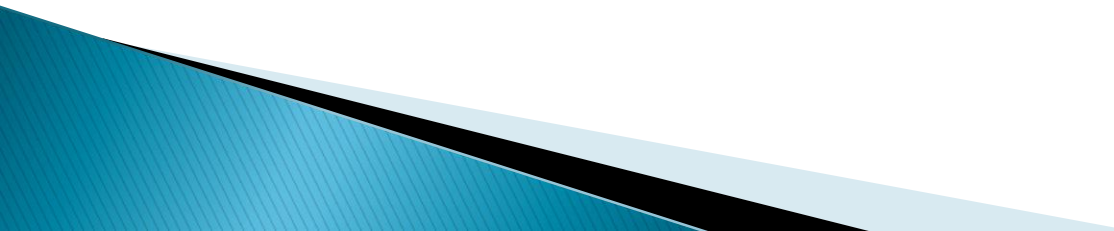
Ilias Tachmazidis, Grigoris Antoniou  
*University of Huddersfield, UK*



# Motivation: The Challenge of Big Data

- ▶ **Big Data:** Huge data set coming from
  - the Web, sensor networks and social media
- ▶ Applications: e.g. smart cities, intelligent environments, information extraction
- ▶ The challenge:
  - Scaling up to big data is not trivial
  - New approaches, new algorithms
- ▶ **Opportunities for Knowledge Representation**
  - Decision making
  - Decision support
  - Data cleaning
  - Inferring high-level knowledge from low-level input

# Outline

- ▶ Motivation
  - ▶ MapReduce paradigm
  - ▶ Logic programming example
    - Joins
    - Anti-joins
  - ▶ The power of stratification
  - ▶ Experimental Results
  - ▶ Future Directions
- 

# The Big Data Challenge for Reasoning

- ▶ Big Data poses significant computational challenges
  - Focus has to be not just complex knowledge structures, but their efficient processing in combination with huge amounts of data
- ▶ In particular for Knowledge Representation (KR), **centralized in-memory solutions** (the traditional KR approach) **do not scale to the Big Data** challenge:
  - Billions of facts result in over 20GB of data

# Related Work

- ▶ **Parallelization approaches:**
  - Rule decomposition
  - Data decomposition
- ▶ **Allows for efficient reasoning on large data sets**
  - **100 billion triples**
  - Datalog (e.g. Afrati & Ullmann)
  - RDF/S (e.g. Weaver & Hendler)
  - OWL dialects (e.g. Urbani et al.)

# Novel Contribution

- ▶ All previous works addressed consistent sets of rules
- ▶ In practice, big data is messy and often inconsistent
- ▶ A type of non-classical reasoning, called **nonmonotonic reasoning**, supports reasoning
  - To deal with inconsistencies that arise naturally in the Web context
  - To deal with deficient (sensor) data
  - To reason with missing (incomplete) information
- ▶ **Apply MapReduce paradigm to nonmonotonic reasoning**

# MapReduce Paradigm

- ▶ Inspired by similar primitives in LISP and other functional languages
- ▶ Operates exclusively on  $\langle \text{key}, \text{value} \rangle$  pairs
- ▶ Input and Output types of a MapReduce job:
  - Input:  $\langle k1, v1 \rangle$
  - $\text{Map}(k1, v1) \rightarrow \text{list}(k2, v2)$
  - $\text{Reduce}(k2, \text{list}(v2)) \rightarrow \text{list}(k3, v3)$
  - Output:  $\text{list}(k3, v3)$

# MapReduce Framework

- ▶ Provides an infrastructure that takes care of
  - distribution of data
  - management of fault tolerance
  - results collection
- ▶ For a specific problem
  - developer writes a few routines which are following the general interface



# Negative Rule Calculation (1 / 5)

- ▶ Models both “join” and “anti-join” operations from database

- ▶ Example:

Facts:

parent(John, Alice), parent(John, Jill),  
sibling(Alice, Edward), sibling(Jill, Mary),  
female(Mary)

Rule:

son(X,Y) ← parent(Y,Z), sibling(Z,X), not female(X)



# Negative Rule Calculation (2 / 5)

## “Join”

### INPUT

Facts in multiple files

#### File01

-----  
parent(John, Alice)  
parent(John, Jill)  
sibling(Alice, Edward)

#### File02

-----  
sibling(Jill, Mary)  
female(Mary)



### MAP phase Input

**Key:** position in file (ignored)

**Value:** fact

<key, parent(John, Alice) >

<key, parent(John, Jill)>

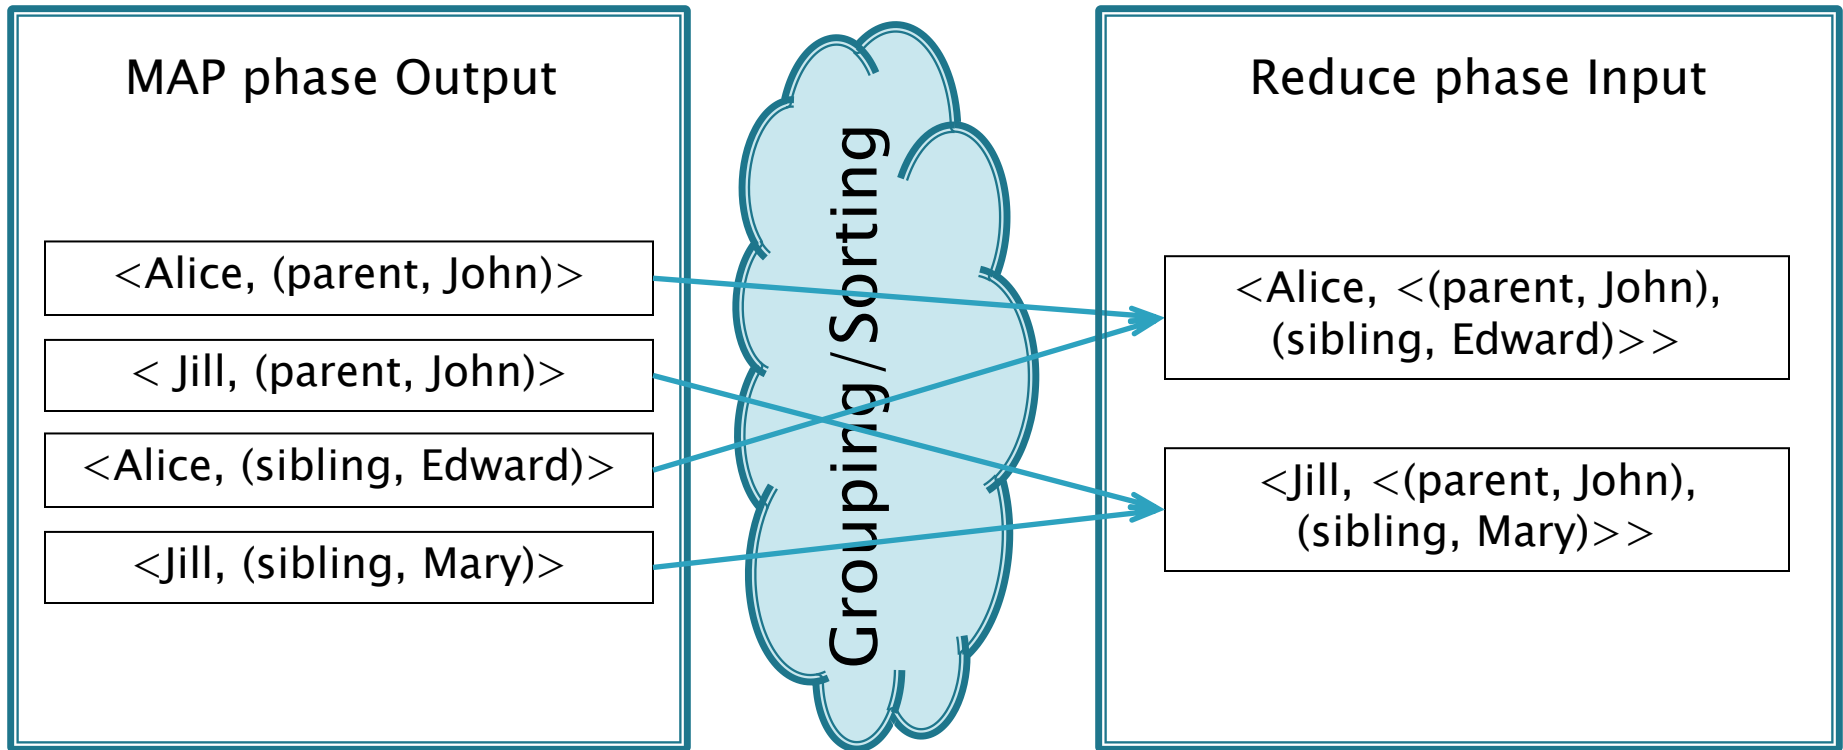
< key, sibling(Alice, Edward)>

<key, sibling(Jill, Mary)>

~~<key, female(Mary)>~~

# Negative Rule Calculation (3 / 5):

## “Join”



# Negative Rule Calculation (4/5)

## “Join”

Reduce phase Input

<Alice, <(parent, John),  
(sibling, Edward)>>

<Jill, <(parent, John),  
(sibling, Mary)>>



Reduce phase Output  
**Output:** new conclusion

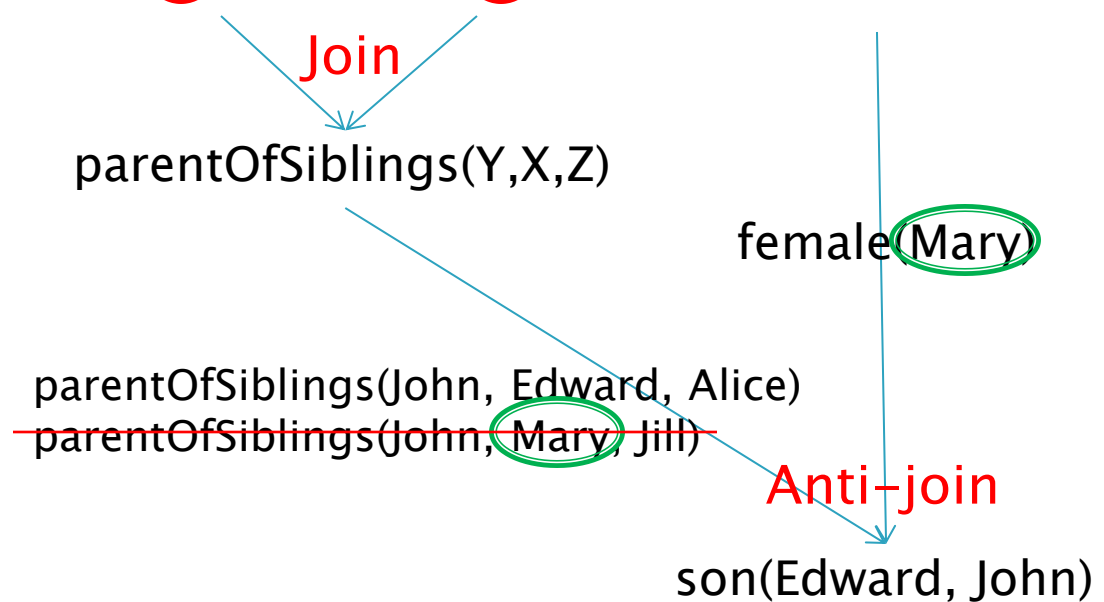
parentOfSiblings(John, Edward,  
Alice)

parentOfSiblings(John, Mary,  
Jill)

# Negative Rule Calculation (5 / 5)

## ▶ Rule:

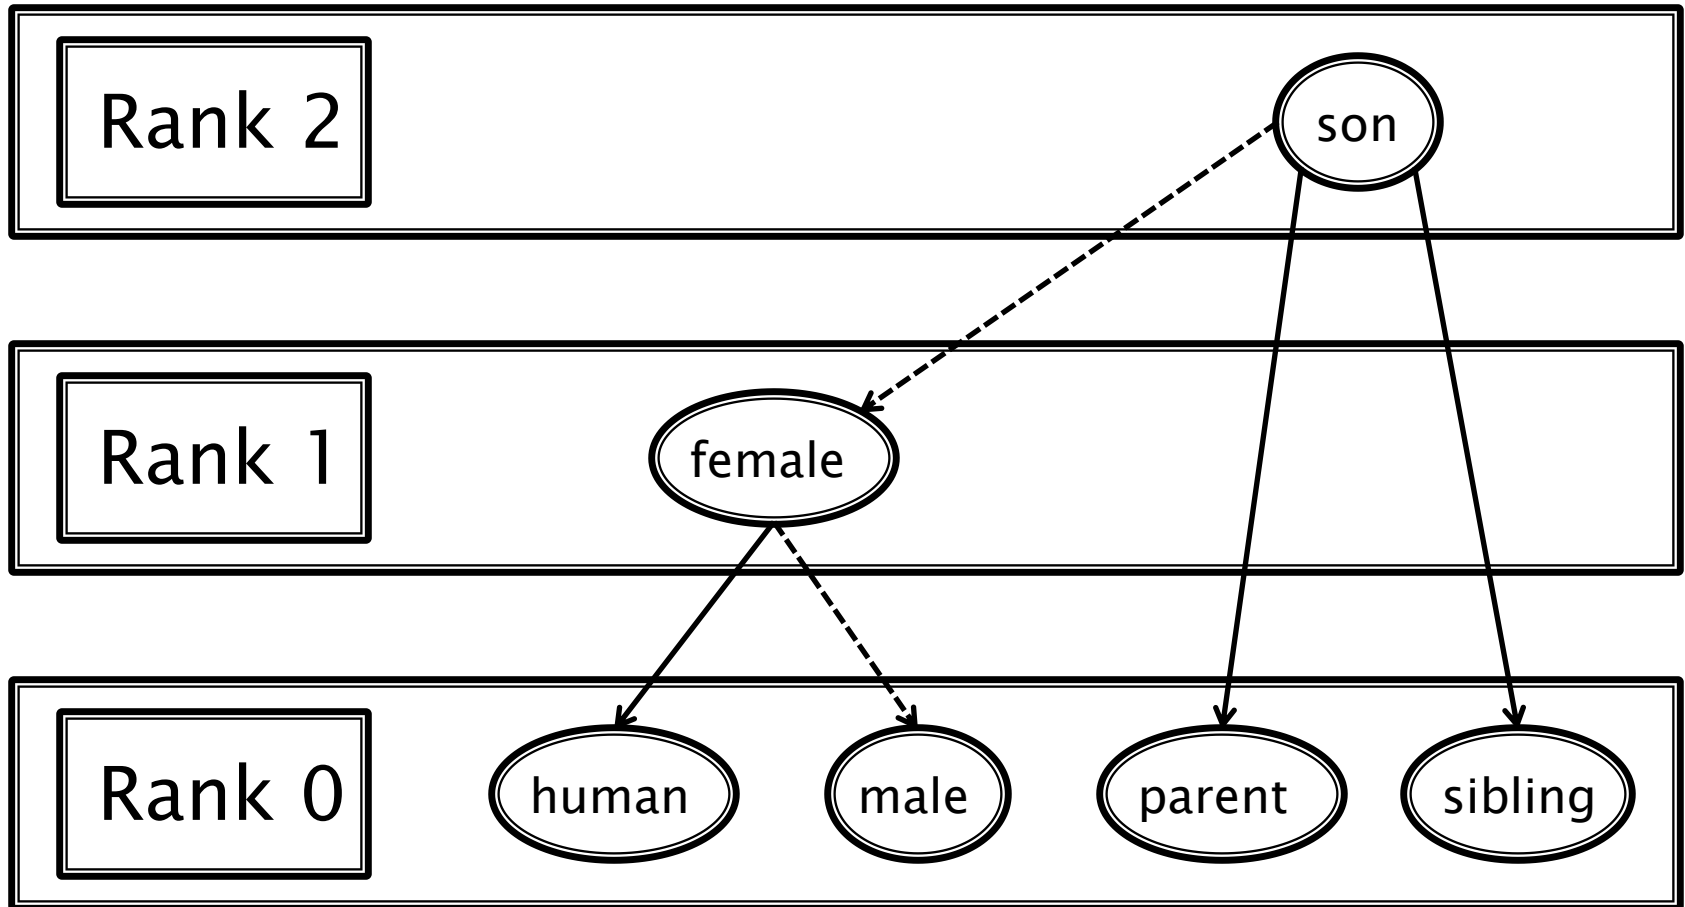
$\text{son}(X,Y) \leftarrow \text{parent}(Y,Z), \text{sibling}(Z,X), \text{not female}(X)$



# Stratified Semantics (1 / 2)

- ▶  $\text{son}(X,Y) \leftarrow \text{parent}(Y,Z), \text{sibling}(Z,X), \text{not female}(X)$
- ▶  $\text{female}(x) \leftarrow \text{human}(x), \text{not male}(x)$

# Stratified Semantics (2/2)



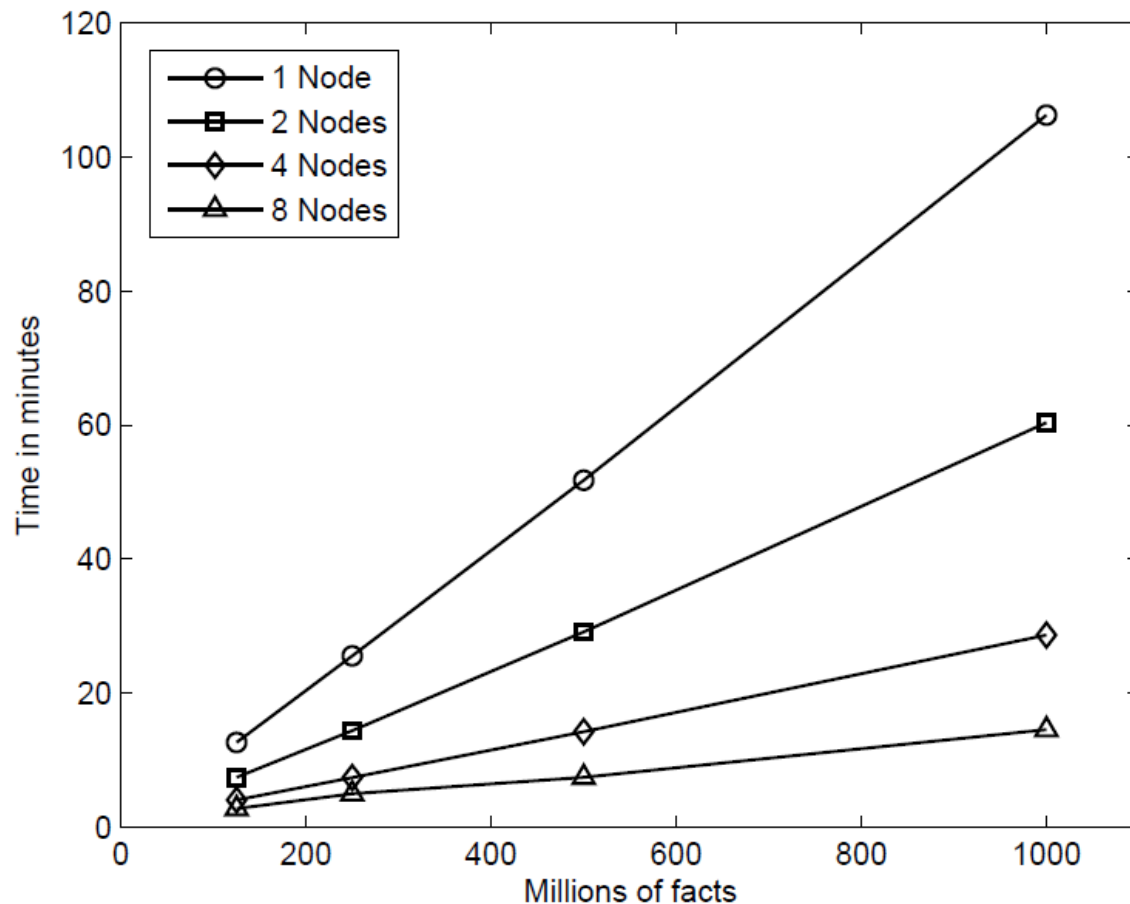
# Experimental Setup

- ▶ Measure **scalability** in terms of
  - Number of nodes (computers in the cluster) maximum parallelization possible
  - Number of facts
  - Number of rules
- ▶ Used a synthetic dataset
  - up to 1 billion facts
  - up to 128 rules



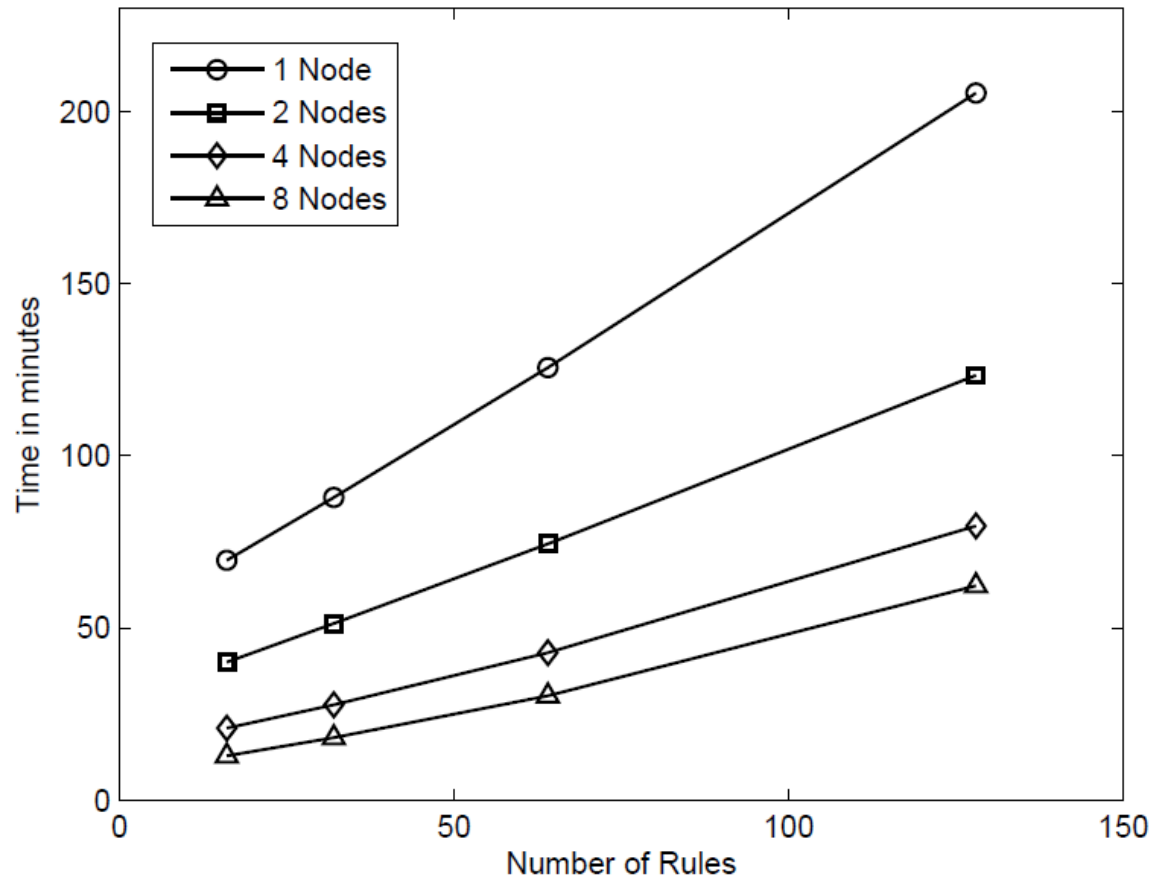
# Experimental Results (1 / 2)

parallelization factor of 8: linear performance



# Experimental Results (2/2)

parallelization factor of 8: linear performance up to 64 rules



# Summary

- ▶ Computed the stratified semantics over Big Data
- ▶ Ran experiments for various
  - data sizes
  - rule sizes
- ▶ Demonstrated that reasoning **can scale well up to 1 billion facts**

# Future Work

## ▶ Beyond stratification

- What happens if we do not have this nice structure
- Solve the problem by allowing dependency cycles

## ▶ Beyond MapReduce

- We will study more complex NMR approaches, including ontology evolution/repair and Answer-Set Programming
- We believe that MapReduce is not well placed to support this kind of approaches: they are probably not “embarrassingly parallel”

# Thank You!



# Experimental Results

